

Machine Learning Project 2025-26

Name:

Kamble Ritesh Dattu (EN24207143)

Kokare Tushar Pandurang (EN24207144)

Kanherkar Amol Sanjay (EN23107056)

Problem Statement :Energy Consumption Forecasting for Smart Homes

```
In [142]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, MinMaxScaler, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import IsolationForest
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [66]: df = pd.read_csv("electricity_bill_dataset.csv")
```

```
In [67]: df
```

Out[67]:

	Fan	Refrigerator	AirConditioner	Television	Monitor	MotorPump	Mont	
0	16	23.0	2.0	6.0	1.0	0	0	0
1	19	22.0	2.0	3.0	1.0	0	0	0
2	7	20.0	2.0	6.0	7.0	0	0	0
3	7	22.0	3.0	21.0	1.0	0	0	0
4	11	23.0	2.0	11.0	1.0	0	0	0
...
45340	18	22.0	3.0	22.0	1.0	0	0	0
45341	23	23.0	2.0	6.0	12.0	0	0	0
45342	22	22.0	2.0	20.0	1.0	0	0	0
45343	8	21.0	2.0	22.0	7.0	0	0	0
45344	8	17.0	2.0	4.0	1.0	0	0	0

45345 rows × 12 columns

Data Pre_proc

In [68]: `df.isnull().sum()`

```
Out[68]: Fan          0  
Refrigerator      0  
AirConditioner    0  
Television        0  
Monitor           0  
MotorPump         0  
Month             0  
City              0  
Company           0  
MonthlyHours      0  
TariffRate        0  
ElectricityBill   0  
dtype: int64
```

```
In [69]: df = df.dropna()  
print(df.isnull().sum())
```

```
Fan          0  
Refrigerator      0  
AirConditioner    0  
Television        0  
Monitor           0  
MotorPump         0  
Month             0  
City              0  
Company           0  
MonthlyHours      0  
TariffRate        0  
ElectricityBill   0  
dtype: int64
```

```
In [70]: df.describe() # Summary statistics for numerical columns
```

	Fan	Refrigerator	AirConditioner	Television	Monitor
count	45345.000000	45345.000000	45345.000000	45345.000000	45345.000000
mean	13.990694	21.705458	1.503959	12.502635	2.865057
std	5.470816	1.672575	1.115482	5.756007	3.894933
min	5.000000	17.000000	0.000000	3.000000	1.000000
25%	9.000000	22.000000	1.000000	7.000000	1.000000
50%	14.000000	22.000000	2.000000	13.000000	1.000000
75%	19.000000	23.000000	2.000000	17.000000	1.000000
max	23.000000	23.000000	3.000000	22.000000	12.000000

```
In [71]: (df.dtypes) # Show data types of each column
```

```
Out[71]: Fan           int64
          Refrigerator   float64
          AirConditioner float64
          Television     float64
          Monitor        float64
          MotorPump      int64
          Month          int64
          City            object
          Company         object
          MonthlyHours    int64
          TariffRate     float64
          ElectricityBill float64
dtype: object
```

```
In [72]: (df.shape) # Output (rows, columns)
```

```
Out[72]: (45345, 12)
```

```
In [73]: (df.sort_index().head())
```

	Fan	Refrigerator	AirConditioner	Television	Monitor	MotorPump	Month			
0	16	23.0	2.0	6.0	1.0	0	10	H		
1	19	22.0	2.0	3.0	1.0	0	5			
2	7	20.0	2.0	6.0	7.0	0	7			
3	7	22.0	3.0	21.0	1.0	0	6			
4	11	23.0	2.0	11.0	1.0	0	2			

Liner Regre

```
In [83]: # Independent Variables (X)
X = df[["Fan", "Refrigerator", "AirConditioner", "Television",
         "Monitor", "MotorPump", "Month", "MonthlyHours", "TariffRate"]]
```

```
In [85]: # Dependent Variable
y = df["ElectricityBill"]
```

```
In [86]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, randc
```

```
In [87]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[87]:
```

LinearRegression

LinearRegression()

```
In [88]: y_pred = model.predict(X_test)
```

```
In [90]: mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Model Evaluation")
print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("R2 Score:", r2)
```

```
Model Evaluation
Mean Squared Error (MSE): 4969.72818804717
Root Mean Squared Error (RMSE): 70.49629910886932
R2 Score: 0.9956383663641158
```

```
In [97]: # Create a DataFrame with Actual vs Predicted values
results = pd.DataFrame({
    "Actual": y_test.values,
    "Predicted": y_pred
})

# Show first 10 rows
print(results.head(10))
```

	Actual	Predicted
0	5054.7	5056.266
1	3809.2	3809.467
2	3245.3	3246.960
3	5728.8	5729.472
4	4148.0	4148.255
5	3645.6	3645.768
6	2317.7	2320.542
7	4100.8	4101.240
8	2286.9	2285.908
9	5394.4	5394.224

```
In [98]: results.to_csv("linear_regression_predictions.csv", index=False)
print("Predictions saved to 'linear_regression_predictions.csv'")
```

```
Predictions saved to 'linear_regression_predictions.csv'
```

Random_fore

```
In [91]: # Step 1 : Split Data into Training & Testing Sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [92]: # Step 2: Build & Train Random Forest Regressor
rf_model = RandomForestRegressor(
    n_estimators=100,          # number of trees
    random_state=42,
    max_depth=None,           # grow trees fully
    min_samples_split=2       # minimum samples to split
)
rf_model.fit(X_train, y_train)
```

```
Out[92]: ▾ RandomForestRegressor ⓘ ?
```

```
RandomForestRegressor(random_state=42)
```


```
In [93]: # Step 3: Make Predictions
y_pred = rf_model.predict(X_test)
```

```
In [94]: # Step 4: Model Evaluation
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```

```
In [95]: print("🎈 Random Forest Model Evaluation")
print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("R² Score:", r2)
```

```
🎈 Random Forest Model Evaluation
Mean Squared Error (MSE): 26.47928567703212
Root Mean Squared Error (RMSE): 5.145802724262962
R² Score: 0.999976760712318
```

```
In [96]: # Actual vs Predicted Bills
plt.figure(figsize=(10,6))
plt.scatter(y_test, y_pred, color='green', alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red')
plt.xlabel("Actual Electricity Bill")
plt.ylabel("Predicted Electricity Bill")
plt.title("Actual vs Predicted Electricity Bill (Random Forest)")
plt.grid(True)
plt.show()
```



```
In [99]: # Create DataFrame for Actual vs Predicted values
rf_results = pd.DataFrame({
    "Actual": y_test.values,
    "Predicted": y_pred
})

# Show first 10 records
print(rf_results.head(10))
```

	Actual	Predicted
0	5054.7	5056.266
1	3809.2	3809.467
2	3245.3	3246.960
3	5728.8	5729.472
4	4148.0	4148.255
5	3645.6	3645.768
6	2317.7	2320.542
7	4100.8	4101.240
8	2286.9	2285.908
9	5394.4	5394.224

```
In [100]: rf_results.to_csv("random_forest_predictions.csv", index=False)
print("Predictions saved to 'random_forest_predictions.csv'")
```

Predictions saved to 'random_forest_predictions.csv'

Isolation Forest for Anomaly Detection

```
In [109...]: # 01) Use only numeric features (exclude City, Company since they are categorical)
features = ["Fan", "Refrigerator", "AirConditioner", "Television",
            "Monitor", "MotorPump", "Month", "MonthlyHours", "TariffRate", "ElectricityBill"]
X = df[features]
```


```
In [110...]: # 02)
iso_forest = IsolationForest(contamination=0.02, random_state=42) # assume 2% anomalies
df["Anomaly"] = iso_forest.fit_predict(X) # -1 = anomaly, 1 = normal
```

```
In [111...]: # 03)
normal_data = df[df["Anomaly"] == 1]
anomalies = df[df["Anomaly"] == -1]

print("Total Records:", len(df))
print("Normal Records:", len(normal_data))
print("Anomalies Detected:", len(anomalies))
```

Total Records: 45345
Normal Records: 44438
Anomalies Detected: 907

```
In [112...]: # 04 ) Plot anomalies vs normal data (ElectricityBill vs MonthlyHours)
plt.figure(figsize=(10,6))
plt.scatter(normal_data["MonthlyHours"], normal_data["ElectricityBill"],
            color='green', label="Normal", alpha=0.5)
plt.scatter(anomalies["MonthlyHours"], anomalies["ElectricityBill"],
            color='red', label="Anomaly", alpha=0.7)
plt.xlabel("Monthly Hours")
plt.ylabel("Electricity Bill")
plt.title("Anomaly Detection in Electricity Bills (Isolation Forest)")
plt.legend()
plt.grid(True)
plt.show()
```



```
In [113]: # Show actual bills with anomaly labels
iso_results = df[["ElectricityBill", "Anomaly"]]
print(iso_results.head(10))
```

	ElectricityBill	Anomaly
0	3225.6	1
1	3806.4	1
2	3203.2	1
3	4370.0	1
4	4204.4	1
5	3485.4	1
6	6417.5	1
7	4182.0	1
8	4641.0	1
9	4392.8	1

```
In [114]: # Normal records
normal_data = df[df["Anomaly"] == 1]

# Anomalous records
anomalies = df[df["Anomaly"] == -1]

print("\nNormal Records Sample:")
print(normal_data.head(5))

print("\nAnomalies Detected Sample:")
print(anomalies.head(5))
```

Normal Records Sample:

	Fan	Refrigerator	AirConditioner	Television	Monitor	MotorPump	Month	\
0	16	23.0	2.0	6.0	1.0	0	10	
1	19	22.0	2.0	3.0	1.0	0	5	
2	7	20.0	2.0	6.0	7.0	0	7	
3	7	22.0	3.0	21.0	1.0	0	6	
4	11	23.0	2.0	11.0	1.0	0	2	

	City	Company	MonthlyHours	\
0	Hyderabad	Tata Power Company Ltd.	384	
1	Vadodara	NHPC	488	
2	Shimla	Jyoti Structure	416	
3	Mumbai	Power Grid Corp	475	
4	Mumbai	Ratnagiri Gas and Power Pvt. Ltd. (RGPPL)	457	

	TariffRate	ElectricityBill	Anomaly
0	8.4	3225.6	1
1	7.8	3806.4	1
2	7.7	3203.2	1
3	9.2	4370.0	1
4	9.2	4204.4	1

Anomalies Detected Sample:

	Fan	Refrigerator	AirConditioner	Television	Monitor	MotorPump	Month	\
171	16	23.0	1.0	22.0	12.0	0	3	
251	22	23.0	0.0	15.0	12.0	0	11	
264	12	23.0	3.0	17.0	12.0	0	12	
277	12	23.0	2.0	20.0	12.0	0	1	
314	17	22.0	3.0	21.0	12.0	0	5	

	City	Company	MonthlyHours	TariffRate	\
171	Navi Mumbai	SJVN Ltd.	768	9.3	
251	Ratnagiri	NLC India	735	7.4	
264	Pune	NLC India	802	9.1	
277	Navi Mumbai	Power Grid Corp	713	9.3	
314	Kolkata	L&T Transmission & Distribution	888	8.7	

	ElectricityBill	Anomaly
171	7142.4	-1
251	5439.0	-1
264	7298.2	-1
277	6630.9	-1
314	7725.6	-1

ARIMA Model

```
In [119]: # Take monthly average electricity bill
ts = df.groupby("Month")["ElectricityBill"].mean()

# Convert numeric months into datetime index (assuming year = 2025)
ts.index = pd.date_range(start="2025-01-01", periods=len(ts), freq="M")
```

```
print(ts)

2025-01-31    4403.964784
2025-02-28    3919.195462
2025-03-31    4409.862445
2025-04-30    4254.603439
2025-05-31    4398.011479
2025-06-30    4219.997220
2025-07-31    4432.699030
2025-08-31    4401.935753
2025-09-30    4221.012469
2025-10-31    4409.195557
2025-11-30    4258.952020
2025-12-31    4422.842174
Freq: ME, Name: ElectricityBill, dtype: float64
C:\Users\kambl\AppData\Local\Temp\ipykernel_20860\1511742778.py:5: FutureWarning:
g: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

ts.index = pd.date_range(start="2025-01-01", periods=len(ts), freq="M")
```

```
In [125...]: # ARIMA order (p,d,q) = (1,1,1) as starting point
model = ARIMA(ts, order=(1,1,1))
model_fit = model.fit()
# Show model summary
print(model_fit.summary())
```

```

SARIMAX Results
=====
Dep. Variable: ElectricityBill No. Observations: 12
Model: ARIMA(1, 1, 1) Log Likelihood -68.906
Date: Fri, 03 Oct 2025 AIC 143.811
Time: 15:41:08 BIC 145.005
Sample: 01-31-2025 HQIC 143.059
- 12-31-2025
Covariance Type: opg
=====
            coef    std err        z      P>|z|      [0.025]    [0.975]
-----
ar.L1     -0.7773    0.167    -4.658      0.000     -1.104    -0.450
ma.L1     -0.3187    0.413    -0.772      0.440     -1.128     0.491
sigma2    1.613e+04  9021.754     1.788      0.074    -1550.708   3.38e+04
=====
=====

Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB):
0.11
Prob(Q): 1.00 Prob(JB):
0.94
Heteroskedasticity (H): 0.26 Skew:
0.05
Prob(H) (two-sided): 0.23 Kurtosis:
2.51
=====

=====
```

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [122]: # Forecast next 6 months
forecast = model_fit.forecast(steps=6)
# Print forecast values
print("Forecasted Electricity Bills:")
print(forecast)
```

```
Forecasted Electricity Bills:
2026-01-31 4280.909265
2026-02-28 4391.229455
2026-03-31 4305.480882
2026-04-30 4372.130674
2026-05-31 4320.325791
2026-06-30 4360.592172
Freq: ME, Name: predicted_mean, dtype: float64
```


```
In [123]: # Plot actual vs forecast
plt.figure(figsize=(10,5))
plt.plot(ts, label="Actual")
plt.plot(pd.date_range(ts.index[-1] + pd.offsets.MonthEnd(1), periods=6, freq='M'),
         forecast, label="Forecast", color="red")
plt.xlabel("Month")
plt.ylabel("Average Electricity Bill")
```

```

plt.title("Electricity Bill Forecast using ARIMA(1,1,1)")
plt.legend()
plt.grid(True)
plt.show()

```

C:\Users\kambl\AppData\Local\Temp\ipykernel_20860\606313886.py:4: FutureWarning:
g: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.
plt.plot(pd.date_range(ts.index[-1] + pd.offsets.MonthEnd(1), periods=6, freq="M")),



In [137]: # Step 2: Predict values for the same period as your dataset
predicted = model_fit.predict(start=0, end=len(ts)-1, typ='levels')

In [138]: # Step 3: Create a comparison DataFrame
comparison = pd.DataFrame({
 'Month': ts.index, # Month or index
 'Actual_Bill': ts.values, # Actual electricity bill
 'Predicted_Bill': predicted # Predicted electricity bill
})


In [139]: # Step 4: Display the comparison
print("Comparison of Actual vs Predicted Electricity Bills:")
print(comparison.head(10)) # Show first 10 rows for example

Comparison of Actual vs Predicted Electricity Bills:

	Month	Actual_Bill	Predicted_Bill
2025-01-31	2025-01-31	4403.964784	0.000000
2025-02-28	2025-02-28	3919.195462	4173.541710
2025-03-31	2025-03-31	4409.862445	4317.027291
2025-04-30	2025-04-30	4254.603439	4000.964228
2025-05-31	2025-05-31	4398.011479	4295.015211
2025-06-30	2025-06-30	4219.997220	4253.742532
2025-07-31	2025-07-31	4432.699030	4369.116465
2025-08-31	2025-08-31	4401.935753	4247.108202
2025-09-30	2025-09-30	4221.012469	4376.502450
2025-10-31	2025-10-31	4409.195557	4411.194537

In [140]: # Step 5: Plot actual vs predicted bills

```
plt.figure(figsize=(10,5))
plt.plot(ts, label="Actual Bills", marker='o')
plt.plot(predicted, label="Predicted Bills", color="green", marker='x')
plt.xlabel("Month")
plt.ylabel("Electricity Bill")
plt.title("Actual vs Predicted Electricity Bills (ARIMA)")
plt.legend()
plt.grid(True)
plt.show()
```



Recommendation Layer (Optimization)

In [152]: # Combine actual and predicted bills

```
recommendation_df = pd.DataFrame({
    'Month': ts.index,
    'ActualBill': ts.values,
    'PredictedBill': predicted # <-- use 'predicted' from ARIMA
```

```
})  
recommendation_df.head()
```

Out[152...]

	Month	ActualBill	PredictedBill
2025-01-31	2025-01-31	4403.964784	0.000000
2025-02-28	2025-02-28	3919.195462	4173.541710
2025-03-31	2025-03-31	4409.862445	4317.027291
2025-04-30	2025-04-30	4254.603439	4000.964228
2025-05-31	2025-05-31	4398.011479	4295.015211

In [156...]

```
# Align lengths  
n = len(ts) # number of months in ts  
subset_df = df.iloc[:n] # take only first n rows  
  
# Create recommendation dataframe  
recommendation_df = pd.DataFrame({  
    'Month': ts.index,  
    'ActualBill': ts.values,  
    'PredictedBill': predicted, # ARIMA predicted values  
    'Fan': subset_df['Fan'].values,  
    'MotorPump': subset_df['MotorPump'].values  
})
```

In [157...]

```
# Generate suggestions  
def generateSuggestions(df):  
    suggestions = []  
    for _, row in df.iterrows():  
        rules = []  
        if row['PredictedBill'] > row['ActualBill'] * 1.2:  
            rules.append("High predicted bill: check appliances")  
        if row['Fan'] > 20:  
            rules.append("Shift Fan usage to off-peak hours")  
        if row['MotorPump'] > 0:  
            rules.append("Check MotorPump for leakage")  
        suggestions.append(", ".join(rules) if rules else "No action needed")  
    return suggestions  
  
recommendation_df['Suggestions'] = generateSuggestions(recommendation_df)  
  
# Display the recommendations  
recommendation_df.head()
```

Out[157...]

	Month	ActualBill	PredictedBill	Fan	MotorPump	Suggestions
2025-01-31	2025-01-31	4403.964784	0.000000	16	0	No action needed
2025-02-28	2025-02-28	3919.195462	4173.541710	19	0	No action needed
2025-03-31	2025-03-31	4409.862445	4317.027291	7	0	No action needed
2025-04-30	2025-04-30	4254.603439	4000.964228	7	0	No action needed
2025-05-31	2025-05-31	4398.011479	4295.015211	11	0	No action needed

In [158...]

```
# Show months where any action is suggested
recommendation_df[recommendation_df['Suggestions'] != "No action needed"]
```

Out[158...]


	Month	ActualBill	PredictedBill	Fan	MotorPump	Suggestions
2025-07-31	2025-07-31	4432.699030	4369.116465	23	0	Shift Fan usage to off peak hours
2025-08-31	2025-08-31	4401.935753	4247.108202	22	0	Shift Fan usage to off peak hours

In [161...]

```
plt.figure(figsize=(12,6))
plt.plot(ts, label='Actual Bills', marker='o')
plt.plot(predicted, label='Predicted Bills', marker='x') # use 'predicted'

# Highlight months needing action
action_months = recommendation_df[recommendation_df['Suggestions'] != "No action needed"]
plt.scatter(action_months, ts[action_months], color='red', s=100, label='Action Months')

plt.xlabel("Month")
plt.ylabel("Electricity Bill")
plt.title("Electricity Bills with Recommendations")
plt.legend()
plt.grid(True)
plt.show()
```



In []: